

Concept Identification Using Lookup Approach

Mohammed Asif Praveen Kumar Senthil Madhappan



Problem Statement

- Design an engine to tag the concepts for the given text using lookup approach. The design should be scalable and update-friendly.
- List of all valid concepts reside on a HBase / NoSQL database and the initial set would be 10 million concepts.
- Benchmark the performance for tagging the document of size 250 words with concepts.

Motivation

Concept identification is becoming an area of great interest and has been extensively used in various applications, notably

- Information Retrieval & Extraction
 - Concept Based Indexing
 - Summarization
 - Question Answering
- Data Mining
 - Document Classification & Categorization
 - Ontology Engineering
 - Literature Based Knowledge Discovery

High-level Solution Approach

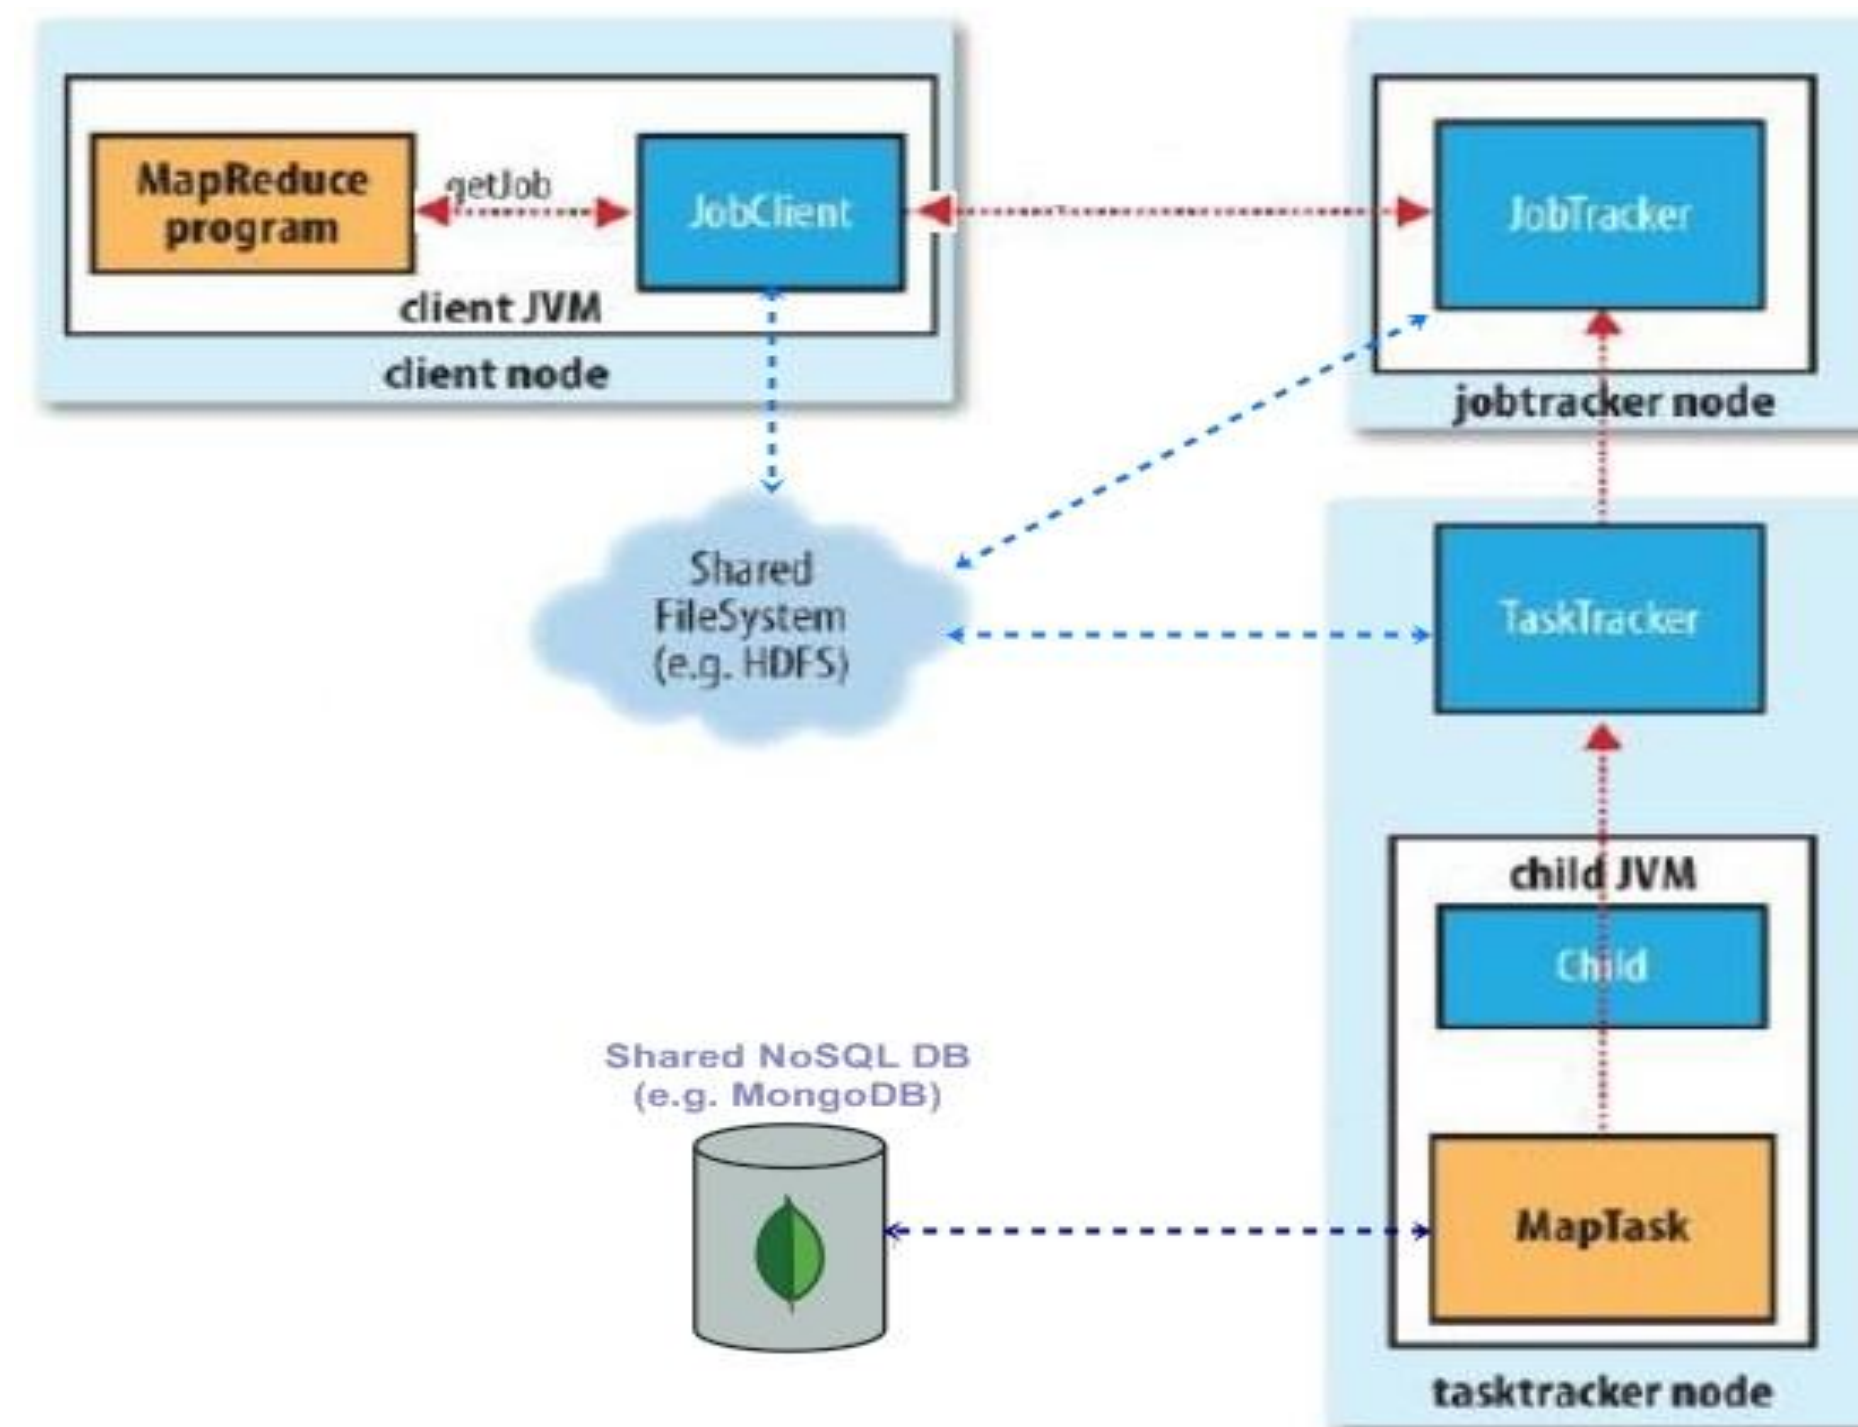
Issues:

- Loading the complete known concepts into memory is expensive in terms of memory foot print that gets added to the JVM heap.
- Lookup cost to check for a concept from known concepts.
- Number of lookups required to find an existing concept.

Solution Approach:

- Store the concepts on a central storage off the JVM heap.
 - The possible options would be HBase (or) NoSQL database like MongoDB
- Time required for the lookup operation can be reduced by
 - Using local/distributed caches between the Hadoop task and the storage.
 - Identifying the words that will never form a concept and avoid lookup for those words e.g. stop words.

Architecture




Performance Tuning

```
> use Cloud
switched to db Cloud
> db.Concepts.count();
8380399
> db.system.indexes.find();
{ "v" : 1, "key" : { "_id" : 1 }, "ns" : "Cloud.Concepts", "name" : "_id_" }
> db.Concepts.find({name:'information technology'}).explain();
{
  "cursor" : "BasicCursor",
  "isMultiKey" : false,
  "n" : 1,
  "nscannedObjects" : 8380399,
  "nscanned" : 8380399,
  "nscannedObjectsAllPlans" : 8380399,
  "nscannedAllPlans" : 8380399,
  "scanAndOrder" : false,
  "indexOnly" : false,
  "nYields" : 20,
  "nChunkSkips" : 0,
  "millis" : 17357,
}
> db.Concepts.ensureIndex({name:1}, {unique:true, dropDups:true});
> db.system.indexes.find();
{ "v" : 1, "key" : { "_id" : 1 }, "ns" : "Cloud.Concepts", "name" : "_id_" }
{ "v" : 1, "key" : { "name" : 1 }, "unique" : true, "ns" : "Cloud.Concepts",
name" : "name_1", "dropDups" : true }
> db.Concepts.find({name:'information technology'}).explain();
{
  "cursor" : "BtreeCursor name_1",
  "isMultiKey" : false,
  "n" : 1,
  "nscannedObjects" : 1,
  "nscanned" : 1,
  "nscannedObjectsAllPlans" : 1,
  "nscannedAllPlans" : 1,
  "scanAndOrder" : false,
  "indexOnly" : false,
  "nYields" : 0,
  "nChunkSkips" : 0,
  "millis" : 5,
}
```

Hardware/Software Configuration

Virtual Machine	Host OS	Windows 7 (64bit)
	Guest OS	Ubuntu 12.04 (64bit)
	RAM	1.5GB
	Processor	2
	Network	Bridged Adapter
Hadoop		Pseudo-Distributed
Mongo DB		Single Centralized Server

Results



Dataset Used:

- Seeded Concepts – 10 Million Wikipedia Titles
- Lookup Documents – 100 Gynanidhi Documents (English)

Future Work

- Experiment with MongoDB replica set configuration
 - Use local database server instead of single centralized server over network.
- Experiment with various database tuning strategies
 - Create additional index and/or change database schema.
 - Tuning database cache parameters.
- Experiment with fully-distributed Hadoop cluster environment.
- Experiment with distributed JVM cache.